

In [25]:

```
import sys
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

# sklearn imports
from sklearn.model_selection import (
    train_test_split,
    GridSearchCV,
    StratifiedKFold,
    cross_val_score
)
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    accuracy_score,
    roc_auc_score,
    roc_curve,
    precision_recall_curve,
    average_precision_score,
    precision_score,
    recall_score,
    f1_score
)

# Advanced models
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier

# Explainability
import shap

# Set random seed for reproducibility
RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)

# =====
# 1. DATA LOADING
# =====
print("">>>> STEP 1: Loading Data...")

file_path = r"C:\Users\admin\Desktop\diabetic_data.csv"

if not os.path.exists(file_path):
    raise FileNotFoundError(f"Error: '{file_path}' not found.")

# Load CSV, treat '?' as missing values
df = pd.read_csv(file_path, na_values="?", low_memory=False)
```

```
# Basic cleaning: remove duplicate rows
df = df.drop_duplicates()

print(f"Data loaded successfully.")
print(f"Shape after dropping duplicates: {df.shape}")
```

>>> STEP 1: Loading Data...
Data loaded successfully.
Shape after dropping duplicates: (101766, 50)

In [26]:

```
# =====
# 1B. EXCLUDE EXPIRED AND HOSPICE PATIENTS
# =====
print("\n>>> STEP 1B: Excluding Expired and Hospice patients...")
exclude_discharge_ids = [11, 13, 14, 19, 20, 21]
before_count = len(df)
df = df[~df['discharge_disposition_id'].isin(exclude_discharge_ids)]
after_count = len(df)
print(f" Excluded {before_count - after_count} samples (Expired/Hospice)")
print(f" Remaining samples: {after_count}")
```

>>> STEP 1B: Excluding Expired and Hospice patients...
Excluded 2423 samples (Expired/Hospice)
Remaining samples: 99343

In [27]:

```
# =====
# 2. FEATURE ENGINEERING
# =====
print("\n>>> STEP 2: Performing Feature Engineering...")

# =====
# 2A. ICD-9 CODE FEATURE EXTRACTION
# =====
diag_cols = ["diag_1", "diag_2", "diag_3"]

def normalize_icd(code):
    """Normalize ICD-9 code into a clean string (None if invalid)."""
    if pd.isna(code):
        return None
    c = str(code).strip()
    if c == "" or c.lower() == "nan":
        return None
    return c

def build_diag_list(row):
    """Build a normalized list of ICD-9 codes for a given row."""
    return [normalize_icd(row[c]) for c in diag_cols]

def has_icd_any(diag_list, icd_set):
    """
    Check if any ICD-9 code in diag_list matches icd_set.
    Match is based on exact code or main part before dot.
    """
    for code in diag_list:
        if code is None:
            continue
        main = code.split(".")[0]
        if code in icd_set or main in icd_set:
            return 1
    return 0
```

```

# === ICD-9 Code Sets (per project instructions) ===
alcohol_codes = {"303", "305", "571", "571.1", "571.2", "571.3"}
bp_codes = {"401", "402", "403", "404", "405", "642"}
cholesterol_codes = {"272"}
heart_codes = {"410", "411", "412", "413", "414", "428"}
obesity_codes = {"278"}
pregnancy_codes = {
    "630", "631", "632", "633", "634", "635", "636", "637", "638", "639",
    "640", "641", "642", "643", "644", "645", "646", "647", "648", "649"
}
uric_acid_codes = {"274", "790.6", "790"}
vision_codes = {"368"}

# Precompute each row's diagnosis list
diag_lists = df.apply(build_diag_list, axis=1)

# Create binary features from ICD-9 codes
df["alcohol"] = diag_lists.apply(lambda dl: has_icd_any(dl, alcohol_codes))
df["blood_pressure"] = diag_lists.apply(lambda dl: has_icd_any(dl, bp_codes))
df["cholesterol"] = diag_lists.apply(lambda dl: has_icd_any(dl, cholesterol_codes))
df["heart_disease"] = diag_lists.apply(lambda dl: has_icd_any(dl, heart_codes))
df["obesity"] = diag_lists.apply(lambda dl: has_icd_any(dl, obesity_codes))
df["pregnancy"] = diag_lists.apply(lambda dl: has_icd_any(dl, pregnancy_codes))
df["uric_acid"] = diag_lists.apply(lambda dl: has_icd_any(dl, uric_acid_codes))
df["blurred_vision"] = diag_lists.apply(lambda dl: has_icd_any(dl, vision_codes))

print("ICD-9 based clinical features created:")
print(f" - alcohol: {df['alcohol'].sum()} positive cases")
print(f" - blood_pressure: {df['blood_pressure'].sum()} positive cases")
print(f" - cholesterol: {df['cholesterol'].sum()} positive cases")
print(f" - heart_disease: {df['heart_disease'].sum()} positive cases")
print(f" - obesity: {df['obesity'].sum()} positive cases")
print(f" - pregnancy: {df['pregnancy'].sum()} positive cases")
print(f" - uric_acid: {df['uric_acid'].sum()} positive cases")
print(f" - blurred_vision: {df['blurred_vision'].sum()} positive cases")

```

>>> STEP 2: Performing Feature Engineering...

ICD-9 based clinical features created:

- alcohol: 2944 positive cases
- blood_pressure: 19667 positive cases
- cholesterol: 2384 positive cases
- heart_disease: 29861 positive cases
- obesity: 1324 positive cases
- pregnancy: 677 positive cases
- uric_acid: 745 positive cases
- blurred_vision: 65 positive cases

In [28]:

```

# =====
# 3. TARGET VARIABLE CREATION (Diabetes Type 2)
# =====
print("\n>>> STEP 3: Creating Target Variable 'diabetes2'...")

def is_type2_diabetes_code(code):
    """
        Check if an ICD-9 code is Type 2 diabetes pattern: 250.x0 or 250.x2.
        Note: 250.x0 is equivalent to 250.x (trailing zero removed)
              250 alone is equivalent to 250.00 (Type 2)
    """
    code = normalize_icd(code)

```

```

if code is None:
    return False

# Case 1: exactly "250" -> equivalent to 250.00 (Type 2)
if code == "250":
    return True

if not code.startswith("250."):
    return False

tail = code.split(".")[1]

# Case 2: 250.x (one digit) -> equivalent to 250.x0 (Type 2)
if len(tail) == 1:
    return True

# Case 3: 250.xy (two digits) -> Type 2 if y is 0 or 2
if len(tail) >= 2:
    return tail[1] in {"0", "2"}

return False

def is_diabetes_250_any(code):
"""
Check if code is 250 or any 250.x pattern.
"""
code = normalize_icd(code)
if code is None:
    return False
return code == "250" or code.startswith("250.")

def define_target(row):
"""
Define Type 2 diabetes target based on project rules:
- Condition 1: At least one diagnosis is 250.x0 or 250.x2.
- OR Condition 2: Diagnosis includes 250 or 250.x, AND
    at least one diabetes medication (metformin, glimepiride, glipizide) is prescribed.

Medication is considered prescribed if value is 'Up', 'Down', or 'Steady' (not 'No')
"""
diags = [row[c] for c in diag_cols]

# Diabetes meds of interest (prescribed if not "No")
meds = [row["metformin"], row["glimepiride"], row["glipizide"]]
meds_prescribed = any(m in ["Up", "Down", "Steady"] for m in meds if pd.notna(m))

# Condition 1: Direct Type 2 code
if any(is_type2_diabetes_code(d) for d in diags):
    return 1

# Condition 2: Any diabetes code + medication prescribed
if meds_prescribed and any(is_diabetes_250_any(d) for d in diags):
    return 1

return 0

df["diabetes2"] = df.apply(define_target, axis=1)

# Summary statistics
total_positive = df["diabetes2"].sum()

```

```

total_negative = len(df) - total_positive
prevalence = total_positive / len(df) * 100

print(f"Target variable 'diabetes2' created.")
print(f" - Total samples: {len(df)}")
print(f" - Positive (1): {total_positive} ({prevalence:.2f}%)")
print(f" - Negative (0): {total_negative} ({100-prevalence:.2f}%)")
print(f" - Class imbalance ratio: 1:{total_negative/total_positive:.2f}")

```

>>> STEP 3: Creating Target Variable 'diabetes2'...

Target variable 'diabetes2' created.

- Total samples: 99343
- Positive (1): 32346 (32.56%)
- Negative (0): 66997 (67.44%)
- Class imbalance ratio: 1:2.07

In [29]:

```

# =====
# 4. DEMOGRAPHICS PROCESSING (Age, Gender, Race)
# =====
print("\n>>> STEP 4: Processing Demographics (Age, Gender, Race)...")

# ---- AGE ----
# Convert age intervals into ordered integer groups 0-9
age_map = {
    "[0-10)": 0, "[10-20)": 1, "[20-30)": 2, "[30-40)": 3, "[40-50)": 4,
    "[50-60)": 5, "[60-70)": 6, "[70-80)": 7, "[80-90)": 8, "[90-100)": 9
}

# Check for invalid age values before filtering
invalid_age = df[~df["age"].isin(age_map.keys())]["age"].unique()
if len(invalid_age) > 0:
    print(f" - Removing {len(df[~df['age'].isin(age_map.keys())])} rows with in")

# Keep only rows where age is valid
df = df[df["age"].isin(age_map.keys())]
df["age"] = df["age"].map(age_map)

print(f" - Age encoded (0-9). Distribution:")
print(df["age"].value_counts().sort_index().to_string())

# ---- GENDER ----
# Remove rows with invalid gender values
valid_genders = {"Male", "Female"}
invalid_gender = df[~df["gender"].isin(valid_genders)]["gender"].unique()
if len(invalid_gender) > 0:
    print(f"\n - Removing {len(df[~df['gender'].isin(valid_genders)])} rows wit")

df = df[df["gender"].isin(valid_genders)]
# Map gender: Male=1, Female=0
df["gender"] = df["gender"].map({"Male": 1, "Female": 0})

print(f"\n - Gender encoded. Distribution:")
print(f"    Male (1): {(df['gender']==1).sum()}")
print(f"    Female (0): {(df['gender']==0).sum()}")

# ---- RACE ----
# Normalize race categories: 1-Caucasian, 2-African American, 3-Asian, 4-Hispanic
def encode_race(x):
    if x == "Caucasian":
        return 1

```

```

    elif x == "AfricanAmerican":
        return 2
    elif x == "Asian":
        return 3
    elif x == "Hispanic":
        return 4
    else:
        # Includes 'Other', missing, unknown values
        return 5

df["race"] = df["race"].apply(encode_race)

print(f"\n - Race encoded. Distribution:")
print(f"  1-Caucasian: {((df['race'] == 1).sum())}")
print(f"  2-African American: {((df['race'] == 2).sum())}")
print(f"  3-Asian: {((df['race'] == 3).sum())}")
print(f"  4-Hispanic: {((df['race'] == 4).sum())}")
print(f"  5-Others: {((df['race'] == 5).sum())}")

```

>>> STEP 4: Processing Demographics (Age, Gender, Race)...
- Age encoded (0-9). Distribution:

age	count
0	160
1	690
2	1649
3	3764
4	9607
5	17060
6	22059
7	25331
8	16434
9	2589

- Removing 3 rows with invalid gender: ['Unknown/Invalid']

- Gender encoded. Distribution:
Male (1): 45886
Female (0): 53454

- Race encoded. Distribution:
1-Caucasian: 74220
2-African American: 18772
3-Asian: 628
4-Hispanic: 2017
5-Others: 3703

In [30]:

```

# =====
# 5. FINAL DATA SELECTION & TRAIN-TEST SPLIT
# =====
print("\n>>> STEP 5: Finalizing Dataset for Modeling...")

# =====
# 5A. DEFINE FEATURE COLUMNS
# =====
# 11 base features (per project requirement) + 3 interaction features
feature_cols = [
    # Demographics (3)
    "race", "gender", "age",
    # ICD-9 based clinical features (8)
    "alcohol", "blood_pressure", "cholesterol", "heart_disease",

```

```

        "obesity", "pregnancy", "uric_acid", "blurred_vision"
    ]

target_col = "diabetes2"

print(f"Features selected ({len(feature_cols)} total):")
print(f" - Demographics (3): race, gender, age")
print(f" - Clinical features (8): alcohol, blood_pressure, cholesterol,")
print(f"   heart_disease, obesity, pregnancy, uric_acid, blurred_vision")
print(f" - Target: {target_col}")

# =====
# 5B. EXTRACT FEATURES AND TARGET
# =====
X = df[feature_cols].copy()
y = df[target_col].astype(int).copy()

# Check for missing values
missing_count = X.isnull().sum().sum()
if missing_count > 0:
    print(f"\n  WARNING: {missing_count} missing values found in features.")
    # Drop rows with missing values
    combined = pd.concat([X, y], axis=1).dropna()
    X = combined[feature_cols]
    y = combined[target_col].astype(int)
    print(f"  Dropped rows with missing values. New shape: {X.shape}")
else:
    print(f"\n  No missing values in features. Shape: {X.shape}")

# =====
# 5C. CLASS BALANCE SUMMARY
# =====
print(f"\nClass distribution in final dataset:")
print(f"  - Class 0 (No diabetes2): {(y==0).sum()} ({(y==0).mean()*100:.2f}%)")
print(f"  - Class 1 (Diabetes2):     {(y==1).sum()} ({(y==1).mean()*100:.2f}%)")

# =====
# 5D. TRAIN-TEST SPLIT
# =====
# 80% train, 20% test, stratified to preserve class proportions
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=RANDOM_SEED,
    stratify=y
)

print(f"\nTrain-Test Split (80/20, stratified):")
print(f"  - Train set: {X_train.shape[0]} samples")
print(f"    - Class 0: {(y_train==0).sum()} ({(y_train==0).mean()*100:.2f}%)")
print(f"    - Class 1: {(y_train==1).sum()} ({(y_train==1).mean()*100:.2f}%)")
print(f"  - Test set: {X_test.shape[0]} samples")
print(f"    - Class 0: {(y_test==0).sum()} ({(y_test==0).mean()*100:.2f}%)")
print(f"    - Class 1: {(y_test==1).sum()} ({(y_test==1).mean()*100:.2f}%)")

# =====
# 5E. FEATURE SUMMARY TABLE
# =====
print("\n" + "="*60)

```

```

print("FEATURE SUMMARY TABLE")
print("*60)
print(f"{'Feature':<25} {'Type':<12} {'Min':>6} {'Max':>6} {'Mean':>8}")
print("-*60)
for col in feature_cols:
    col_type = "Binary" if X[col].nunique() <= 2 else "Ordinal"
    print(f"{col:<25} {col_type:<12} {X[col].min():>6.0f} {X[col].max():>6.0f} {X[col].mean():>8.0f}")
print("*60)

```

>>> STEP 5: Finalizing Dataset for Modeling...

Features selected (11 total):

- Demographics (3): race, gender, age
- Clinical features (8): alcohol, blood_pressure, cholesterol, heart_disease, obesity, pregnancy, uric_acid, blurred_vision
- Target: diabetes2

No missing values in features. Shape: (99340, 11)

Class distribution in final dataset:

- Class 0 (No diabetes2): 66994 (67.44%)
- Class 1 (Diabetes2): 32346 (32.56%)

Train-Test Split (80/20, stratified):

- Train set: 79472 samples
 - Class 0: 53595 (67.44%)
 - Class 1: 25877 (32.56%)
- Test set: 19868 samples
 - Class 0: 13399 (67.44%)
 - Class 1: 6469 (32.56%)

FEATURE SUMMARY TABLE					
Feature	Type	Min	Max	Mean	
race	Ordinal	1	5	1.4116	
gender	Binary	0	1	0.4619	
age	Ordinal	0	9	6.0744	
alcohol	Binary	0	1	0.0296	
blood_pressure	Binary	0	1	0.1980	
cholesterol	Binary	0	1	0.0240	
heart_disease	Binary	0	1	0.3006	
obesity	Binary	0	1	0.0133	
pregnancy	Binary	0	1	0.0068	
uric_acid	Binary	0	1	0.0075	
blurred_vision	Binary	0	1	0.0007	

In [31]:

```

# =====
# 5F. DESCRIPTIVE ANALYSIS
# =====
print("\n" + "*80)
print(">>> STEP 5F: DESCRIPTIVE ANALYSIS OF FEATURES")
print("*80)
print("(This section demonstrates data preprocessing was done correctly)")

# -----
# 5F-1. Basic Information of Sample
# -----
print("\n" + "-*60)

```

```

print("5F-1. SAMPLE CHARACTERISTICS")
print("-"*60)
print(f"Total samples: {len(df)}")
print(f"Total features: {len(feature_cols)}")
print(f"Target variable: diabetes2")
print(f" - Positive (Type 2 Diabetes): {(y==1).sum()} ({(y==1).mean() * 100:.2f}%)")
print(f" - Negative (No Diabetes): {(y==0).sum()} ({(y==0).mean() * 100:.2f}%)")
print(f" - Imbalance Ratio: 1:{(y==0).sum()}/{(y==1).sum():.2f}")

# -----
# 5F-2. Distribution of Demographic Characteristics
# -----
print("\n" + "-"*60)
print("5F-2. DEMOGRAPHIC FEATURES DISTRIBUTION")
print("-"*60)

# Age distribution
print("\n[Age Distribution]")
age_dist = df.groupby('age').agg({
    'diabetes2': ['count', 'sum', 'mean']
}).round(4)
age_dist.columns = ['Total', 'Diabetes_Count', 'Diabetes_Rate']
age_labels = ['[0-10)', '[10-20)', '[20-30)', '[30-40)', '[40-50)',
              '[50-60)', '[60-70)', '[70-80)', '[80-90)', '[90-100)']
age_dist.index = [age_labels[i] for i in age_dist.index]
print(age_dist.to_string())

# Gender distribution
print("\n[Gender Distribution]")
gender_dist = df.groupby('gender').agg({
    'diabetes2': ['count', 'sum', 'mean']
}).round(4)
gender_dist.columns = ['Total', 'Diabetes_Count', 'Diabetes_Rate']
gender_dist.index = ['Female', 'Male']
print(gender_dist.to_string())

# Race distribution
print("\n[Race Distribution]")
race_dist = df.groupby('race').agg({
    'diabetes2': ['count', 'sum', 'mean']
}).round(4)
race_dist.columns = ['Total', 'Diabetes_Count', 'Diabetes_Rate']
race_labels = {1: 'Caucasian', 2: 'African American', 3: 'Asian', 4: 'Hispanic'}
race_dist.index = [race_labels[i] for i in race_dist.index]
print(race_dist.to_string())

# -----
# 5F-3. Distribution of Clinical Features (ICD-9 based)
# -----
print("\n" + "-"*60)
print("5F-3. CLINICAL FEATURES DISTRIBUTION (ICD-9 Based)")
print("-"*60)

clinical_features = ['alcohol', 'blood_pressure', 'cholesterol', 'heart_disease',
                     'obesity', 'pregnancy', 'uric_acid', 'blurred_vision']

clinical_summary = []
for feat in clinical_features:
    total_positive = df[feat].sum()
    diabetes_with_feat = df[df[feat]==1]['diabetes2'].sum()

```

```

diabetes_rate_with = df[df[feat]==1]['diabetes2'].mean() if total_positive >
diabetes_rate_without = df[df[feat]==0]['diabetes2'].mean()

clinical_summary.append({
    'Feature': feat,
    'Positive_Cases': total_positive,
    'Prevalence(%)': total_positive/len(df)*100,
    'Diabetes_Rate_With(%)': diabetes_rate_with*100,
    'Diabetes_Rate_Without(%)': diabetes_rate_without*100,
    'Relative_Risk': diabetes_rate_with/diabetes_rate_without if diabetes_ra
})

clinical_df = pd.DataFrame(clinical_summary)
print(clinical_df.to_string(index=False))

# -----
# 5F-4. The Correlation Between Features and Target Variables
# -----
print("\n" + "-"*60)
print("5F-4. FEATURE-TARGET CORRELATION")
print("-"*60)

correlations = X.corrwith(y).sort_values(ascending=False)
print("\nCorrelation with diabetes2 (sorted):")
for feat, corr in correlations.items():
    print(f" {feat}: {corr:.2f}")

# -----
# 5F-5. Descriptive Statistical Visualization
# -----
print("\n>>> Generating Descriptive Analysis Figures...")

# Figure 1: Distribution of target variables (pie chart)
fig, axes = plt.subplots(1, 3, figsize=(15, 5))

# 1a. Target variable distribution
labels = ['Non-Diabetic', 'Type 2 Diabetic']
sizes = [(y==0).sum(), (y==1).sum()]
colors = ['#66b3ff', '#ff6666']
axes[0].pie(sizes, labels=labels, colors=colors, autopct='%.1f%%', startangle=90)
axes[0].set_title('Target Variable Distribution\n(diabetes2)', fontsize=12)

# 1b. Age distribution with diabetes rate
age_counts = df.groupby('age').size()
diabetes_rates = df.groupby('age')['diabetes2'].mean() * 100
ax1b = axes[1]
ax1b_twin = ax1b.twinx()
bars = ax1b.bar(range(len(age_counts)), age_counts.values, color='steelblue', align='center')
line = ax1b_twin.plot(range(len(age_counts)), diabetes_rates.values, 'r-o', linewidth=2)
ax1b.set_xlabel('Age Group')
ax1b.set_ylabel('Count', color='steelblue')
ax1b_twin.set_ylabel('Diabetes Rate (%)', color='red')
ax1b.set_xticks(range(len(age_counts)))
ax1b.set_xticklabels(['0-10', '10-20', '20-30', '30-40', '40-50', '50-60', '60-70'])
ax1b.set_title('Age Distribution & Diabetes Rate', fontsize=12)

# 1c. Clinical features prevalence
feat_names = ['alcohol', 'blood_pr', 'cholest', 'heart_dis', 'obesity', 'pregn', 'smoke']
feat_prevalence = [df[f].mean()*100 for f in clinical_features]
axes[2].barh(feat_names, feat_prevalence, color='teal')

```

```

axes[2].set_xlabel('Prevalence (%)')
axes[2].set_title('Clinical Features Prevalence', fontsize=12)
for i, v in enumerate(feat_prevalence):
    axes[2].text(v + 0.5, i, f'{v:.1f}%', va='center')

plt.tight_layout()
plt.savefig('descriptive_analysis_1.png', dpi=300, bbox_inches='tight')
plt.show()
print("Saved: descriptive_analysis_1.png")

# Figure 2: Feature correlation heat map
plt.figure(figsize=(12, 10))
correlation_matrix = X.corr()
import seaborn as sns
sns.heatmap(correlation_matrix, annot=True, cmap='RdBu_r', center=0,
             fmt='.2f', square=True, linewidths=0.5)
plt.title('Feature Correlation Matrix', fontsize=14)
plt.tight_layout()
plt.savefig('correlation_heatmap.png', dpi=300, bbox_inches='tight')
plt.show()
print("Saved: correlation_heatmap.png")

# Figure 3: Comparison of diabetes rates for each clinical feature
fig, axes = plt.subplots(2, 4, figsize=(16, 8))
axes = axes.flatten()

for i, feat in enumerate(clinical_features):
    ax = axes[i]
    # The rate of diabetes with and without this feature
    rates = [
        df[df[feat]==0]['diabetes2'].mean() * 100,
        df[df[feat]==1]['diabetes2'].mean() * 100
    ]
    bars = ax.bar(['Without', 'With'], rates, color=[ '#3498db', '#e74c3c' ])
    ax.set_title(feat.replace('_', ' ').title(), fontsize=11)
    ax.set_ylabel('Diabetes Rate (%)')
    ax.set_ylim(0, max(rates)*1.3 if max(rates) > 0 else 30)
    for bar, rate in zip(bars, rates):
        ax.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
                f'{rate:.1f}%', ha='center', fontsize=9)

plt.suptitle('Diabetes Rate by Clinical Feature Presence', fontsize=14)
plt.tight_layout()
plt.savefig('clinical_features_diabetes_rate.png', dpi=300, bbox_inches='tight')
plt.show()
print("Saved: clinical_features_diabetes_rate.png")

print("\n>>> Descriptive Analysis complete.")

```

>>> STEP 5F: DESCRIPTIVE ANALYSIS OF FEATURES

(This section demonstrates data preprocessing was done correctly)

5F-1. SAMPLE CHARACTERISTICS

Total samples: 99340

Total features: 11

Target variable: diabetes2

- Positive (Type 2 Diabetes): 32346 (32.56%)
- Negative (No Diabetes): 66994 (67.44%)
- Imbalance Ratio: 1:2.07

5F-2. DEMOGRAPHIC FEATURES DISTRIBUTION

[Age Distribution]

	Total	Diabetes_Count	Diabetes_Rate
[0-10)	160	8	0.0500
[10-20)	690	136	0.1971
[20-30)	1649	577	0.3499
[30-40)	3764	1720	0.4570
[40-50)	9607	4297	0.4473
[50-60)	17060	6746	0.3954
[60-70)	22058	7265	0.3294
[70-80)	25329	7143	0.2820
[80-90)	16434	3888	0.2366
[90-100)	2589	566	0.2186

[Gender Distribution]

	Total	Diabetes_Count	Diabetes_Rate
Female	53454	17284	0.3233
Male	45886	15062	0.3282

[Race Distribution]

	Total	Diabetes_Count	Diabetes_Rate
Caucasian	74220	22881	0.3083
African American	18772	7060	0.3761
Asian	628	202	0.3217
Hispanic	2017	844	0.4184
Others	3703	1359	0.3670

5F-3. CLINICAL FEATURES DISTRIBUTION (ICD-9 Based)

	Feature	Positive_Cases	Prevalence(%)	Diabetes_Rate_With(%)	Diabetes_Rate
	te_Without(%)	Relative_Risk			
32.798041	alcohol	2944	2.963559	24.796196	
31.414263	blood_pressure	19666	19.796658	37.206346	
32.319815	cholesterol	2384	2.399839	42.365772	
37.428037	heart_disease	29860	30.058385	21.235767	
32.332476	obesity	1324	1.332796	49.471299	
		1.530081			


pregnancy	677	0.681498	29.837518
32.579589	0.915835		
uric_acid	745	0.749950	20.000000
32.655814	0.612448		
blurred_vision	65	0.065432	21.538462
32.568119	0.661336		

5F-4. FEATURE-TARGET CORRELATION


Correlation with diabetes2 (sorted):

race	:	0.0501
blood_pressure	:	0.0493
obesity	:	0.0419
cholesterol	:	0.0328
gender	:	0.0052
pregnancy	:	-0.0048
blurred_vision	:	-0.0060
uric_acid	:	-0.0233
alcohol	:	-0.0290
age	:	-0.1256
heart_disease	:	-0.1584


>>> Generating Descriptive Analysis Figures...



Saved: descriptive_analysis_1.png



Saved: correlation_heatmap.png



Saved: clinical_features_diabetes_rate.png

>>> Descriptive Analysis complete.

```
In [32]: # =====
# 6. MODEL TRAINING AND EVALUATION
# =====
print("\n>>> STEP 6: Model Training...")
# =====
```

```

# 6A. DEFINE EVALUATION HELPER FUNCTION & CV STRATEGIES
# =====

# Define cross-validation strategies (used consistently across all models)
cv_5fold = StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_SEED)
cv_10fold = StratifiedKFold(n_splits=10, shuffle=True, random_state=RANDOM_SEED)

print("Cross-validation strategies defined:")
print(" - 5-fold: for hyperparameter tuning (GridSearchCV)")
print(" - 10-fold: for final model evaluation (paper-style)")

def evaluate_model(name, y_true, y_pred, y_proba, print_results=True):
    """
    Comprehensive model evaluation function.
    Returns a dictionary with all metrics for later comparison.
    """

    # Confusion Matrix components
    cm = confusion_matrix(y_true, y_pred)
    tn, fp, fn, tp = cm.ravel()

    # Calculate metrics
    acc = accuracy_score(y_true, y_pred)
    prec = precision_score(y_true, y_pred, zero_division=0)
    rec = recall_score(y_true, y_pred, zero_division=0)
    f1 = f1_score(y_true, y_pred, zero_division=0)
    roc_auc = roc_auc_score(y_true, y_proba)
    pr_auc = average_precision_score(y_true, y_proba)

    # Specificity (True Negative Rate)
    specificity = tn / (tn + fp) if (tn + fp) > 0 else 0

    # Store results
    results = {
        "name": name,
        "accuracy": acc,
        "precision": prec,
        "recall": rec,
        "f1_score": f1,
        "specificity": specificity,
        "roc_auc": roc_auc,
        "pr_auc": pr_auc,
        "tn": tn, "fp": fp, "fn": fn, "tp": tp,
        "y_pred": y_pred,
        "y_proba": y_proba
    }

    if print_results:
        print(f"\n{'='*60}")
        print(f" {name}")
        print(f"{'='*60}")
        print(f"\nConfusion Matrix:")
        print(f" Predicted")
        print(f" Neg Pos")
        print(f" Actual Neg {tn:>5} {fp:>5} (TN, FP)")
        print(f" Actual Pos {fn:>5} {tp:>5} (FN, TP)")
        print(f"\nMetrics:")
        print(f" Accuracy: {acc:.4f}")
        print(f" Precision: {prec:.4f}")
        print(f" Recall: {rec:.4f}")
        print(f" F1-Score: {f1:.4f}")

```

```

        print(f" Specificity: {specificity:.4f}")
        print(f" ROC AUC:      {roc_auc:.4f}")
        print(f" PR AUC:       {pr_auc:.4f}")

    return results

# Dictionary to store all model results (test set evaluation)
all_results = {}

# Dictionary to store trained models (for SHAP analysis later)
trained_models = {}

print("\nEvaluation function defined. Ready for model training.")

# =====
# 6A-2. FEATURE SELECTION ANALYSIS
# =====
print("\n" + "="*60)
print(">>> STEP 6A-2: FEATURE SELECTION ANALYSIS")
print("=".*60)

from sklearn.feature_selection import SelectKBest, mutual_info_classif, chi2

# method1: Information Gain (Mutual Information)
print("\n--- Method 1: Information Gain (Mutual Information) ---")
mi_selector = SelectKBest(score_func=mutual_info_classif, k='all')
mi_selector.fit(X_train, y_train)
mi_scores = pd.Series(mi_selector.scores_, index=feature_cols).sort_values(ascending=False)

# method2: Chi-squared Test
print("--- Method 2: Chi-squared Test ---")
chi2_selector = SelectKBest(score_func=chi2, k='all')
chi2_selector.fit(X_train, y_train)
chi2_scores = pd.Series(chi2_selector.scores_, index=feature_cols).sort_values(ascending=False)

# method3: Random Forest Feature Importance
print("--- Method 3: Random Forest Importance ---")
rf_for_fs = RandomForestClassifier(n_estimators=100, class_weight='balanced',
                                   random_state=RANDOM_SEED, n_jobs=-1)
rf_for_fs.fit(X_train, y_train)
rf_importance = pd.Series(rf_for_fs.feature_importances_, index=feature_cols).sort_values(ascending=False)

# method4: Correlation with Target
print("--- Method 4: Correlation with Target ---")
corr_scores = X_train.corrwith(y_train).abs().sort_values(ascending=False)

# Summary table
print("\n" + "="*80)
print("FEATURE SELECTION RESULTS SUMMARY")
print("=".*80)

fs_summary = pd.DataFrame({
    'Feature': feature_cols,
    'MI_Score': mi_scores.reindex(feature_cols).values,
    'Chi2_Score': chi2_scores.reindex(feature_cols).values,
    'RF_Importance': rf_importance.reindex(feature_cols).values,
    'Correlation': corr_scores.reindex(feature_cols).values
})

# Calculate the average ranking

```

```

for col in ['MI_Score', 'Chi2_Score', 'RF_Importance', 'Correlation']:
    fs_summary[f'{col}_Rank'] = fs_summary[col].rank(ascending=False)

fs_summary['Avg_Rank'] = fs_summary[['MI_Score_Rank', 'Chi2_Score_Rank',
                                      'RF_Importance_Rank', 'Correlation_Rank']]
fs_summary = fs_summary.sort_values('Avg_Rank')

print(fs_summary[['Feature', 'MI_Score', 'Chi2_Score', 'RF_Importance', 'Correla

# Visualization
fig, axes = plt.subplots(1, 4, figsize=(20, 5))

# MI Scores
axes[0].barh(mi_scores.index[::-1], mi_scores.values[::-1], color='steelblue')
axes[0].set_title('Information Gain', fontsize=12)
axes[0].set_xlabel('MI Score')

# Chi2 Scores
axes[1].barh(chi2_scores.index[::-1], chi2_scores.values[::-1], color='coral')
axes[1].set_title('Chi-squared Score', fontsize=12)
axes[1].set_xlabel('Chi2 Score')

# RF Importance
axes[2].barh(rf_importance.index[::-1], rf_importance.values[::-1], color='forestgreen')
axes[2].set_title('Random Forest Importance', fontsize=12)
axes[2].set_xlabel('Importance')

# Correlation
axes[3].barh(corr_scores.index[::-1], corr_scores.values[::-1], color='purple')
axes[3].set_title('Correlation with Target', fontsize=12)
axes[3].set_xlabel('|Correlation|')

plt.suptitle('Feature Selection Methods Comparison', fontsize=14)
plt.tight_layout()
plt.savefig('feature_selection_comparison.png', dpi=300, bbox_inches='tight')
plt.show()
print("\nSaved: feature_selection_comparison.png")

print("\n>>> Feature Selection Analysis complete.")

# =====
# 6B. BASELINE MODEL: Random Forest
# =====
print("\n" + "*60")
print(">>> STEP 6B: Training Baseline Random Forest...")
print("*60")

rf_baseline = RandomForestClassifier(
    n_estimators=300,
    max_depth=10,
    min_samples_leaf=2,
    class_weight="balanced",
    random_state=RANDOM_SEED,
    n_jobs=-1
)
rf_baseline.fit(X_train, y_train)

y_pred_rf_base = rf_baseline.predict(X_test)

```

```

y_proba_rf_base = rf_baseline.predict_proba(X_test)[:, 1]

all_results["RF_Baseline"] = evaluate_model(
    "Baseline Random Forest",
    y_test, y_pred_rf_base, y_proba_rf_base
)

trained_models["RF_Baseline"] = rf_baseline

print("\n>>> Baseline Random Forest complete.")

```

>>> STEP 6: Model Training...

Cross-validation strategies defined:

- 5-fold: for hyperparameter tuning (GridSearchCV)
- 10-fold: for final model evaluation (paper-style)

Evaluation function defined. Ready for model training.

```
=====
>>> STEP 6A-2: FEATURE SELECTION ANALYSIS
=====
```

--- Method 1: Information Gain (Mutual Information) ---


--- Method 2: Chi-squared Test ---

--- Method 3: Random Forest Importance ---

--- Method 4: Correlation with Target ---

```
=====
FEATURE SELECTION RESULTS SUMMARY
=====
```

Feature	MI_Score	Chi2_Score	RF_Importance	Correlation	Avg_Rank
heart_disease	0.016989	1441.934197	0.341047	0.161145	1.250
age	0.012531	538.385411	0.341658	0.127417	1.750
race	0.006488	124.499412	0.086384	0.051830	3.750
blood_pressure	0.000000	154.531010	0.066313	0.049242	4.875
cholesterol	0.000000	74.746019	0.053921	0.031042	6.375
obesity	0.000000	141.816102	0.021863	0.042518	6.375
alcohol	0.000000	66.670977	0.034350	0.029406	7.125
uric_acid	0.000089	42.882801	0.016804	0.023315	7.500
gender	0.006650	0.043015	0.025240	0.001003	8.000
pregnancy	0.000000	2.604481	0.008125	0.005744	9.125
blurred_vision	0.000000	2.281788	0.004296	0.005360	9.875



Saved: feature_selection_comparison.png

>>> Feature Selection Analysis complete.

=====

>>> STEP 6B: Training Baseline Random Forest...

=====

=====

Baseline Random Forest

=====

Confusion Matrix:

	Predicted		
	Neg	Pos	
Actual Neg	7981	5418	(TN, FP)
Actual Pos	2548	3921	(FN, TP)

Metrics:

Accuracy: 0.5991
 Precision: 0.4199
 Recall: 0.6061
 F1-Score: 0.4961
 Specificity: 0.5956
 ROC AUC: 0.6418
 PR AUC: 0.4496

>>> Baseline Random Forest complete.

```
In [33]: # =====
# 6C. TUNED MODEL: Random Forest with GridSearchCV
# =====
print("\n" + "="*60)
print(">>> STEP 6C: Tuning Random Forest (GridSearchCV, 5-fold)...")
print("=".*60)

param_grid_rf = {
    "n_estimators": [100, 200, 300],
    "max_depth": [6, 8, 10, 12],
    "min_samples_split": [2, 5],
    "min_samples_leaf": [1, 2, 4]
}

rf_for_tuning = RandomForestClassifier(
    class_weight="balanced",
    random_state=RANDOM_SEED,
    n_jobs=-1
)

grid_search_rf = GridSearchCV(
    estimator=rf_for_tuning,
    param_grid=param_grid_rf,
    scoring="f1",
    cv=cv_5fold,
    n_jobs=-1,
    verbose=1
)

grid_search_rf.fit(X_train, y_train)
```

```

print(f"\nBest parameters: {grid_search_rf.best_params_}")
print(f"Best CV F1 score: {grid_search_rf.best_score_:.4f}")

rf_tuned = grid_search_rf.best_estimator_

y_pred_rf_tuned = rf_tuned.predict(X_test)
y_proba_rf_tuned = rf_tuned.predict_proba(X_test)[:, 1]

all_results["RF_Tuned"] = evaluate_model(
    "Tuned Random Forest",
    y_test, y_pred_rf_tuned, y_proba_rf_tuned
)

trained_models["RF_Tuned"] = rf_tuned

print("\n>>> Tuned Random Forest complete.")

```

```

=====
>>> STEP 6C: Tuning Random Forest (GridSearchCV, 5-fold)...
=====
Fitting 5 folds for each of 72 candidates, totalling 360 fits

Best parameters: {'max_depth': 6, 'min_samples_leaf': 4, 'min_samples_split': 2,
'n_estimators': 100}
Best CV F1 score: 0.5039

```

```

=====
Tuned Random Forest
=====
```

Confusion Matrix:

		Predicted	
		Neg	Pos
Actual	Neg	7806	5593 (TN, FP)
	Pos	2480	3989 (FN, TP)

Metrics:

Accuracy:	0.5937
Precision:	0.4163
Recall:	0.6166
F1-Score:	0.4970
Specificity:	0.5826
ROC AUC:	0.6377
PR AUC:	0.4451

>>> Tuned Random Forest complete.

```

In [34]: # =====
# 6D. ADVANCED MODEL: XGBoost
# =====
print("\n" + "="*60)
print(">>> STEP 6D: Training XGBoost...")
print("=".*60)

# Compute scale_pos_weight to handle class imbalance
scale_pos_weight = (y_train == 0).sum() / (y_train == 1).sum()
print(f"Class imbalance ratio (scale_pos_weight): {scale_pos_weight:.2f}")

xgb_model = XGBClassifier(
    n_estimators=200,

```

```

        max_depth=6,
        learning_rate=0.1,
        subsample=0.8,
        colsample_bytree=0.8,
        objective="binary:logistic",
        eval_metric="logloss",
        scale_pos_weight=scale_pos_weight,
        random_state=RANDOM_SEED,
        n_jobs=-1,
        verbosity=0
    )

xgb_model.fit(X_train, y_train)

y_pred_xgb = xgb_model.predict(X_test)
y_proba_xgb = xgb_model.predict_proba(X_test)[:, 1]

all_results["XGBoost"] = evaluate_model(
    "XGBoost",
    y_test, y_pred_xgb, y_proba_xgb
)

trained_models["XGBoost"] = xgb_model

print("\n>>> XGBoost complete.")

```

```
=====
>>> STEP 6D: Training XGBoost...
=====
```

```
Class imbalance ratio (scale_pos_weight): 2.07
```

```
=====
XGBoost
=====
```

Confusion Matrix:

		Predicted	
		Neg	Pos
Actual	Neg	8027	5372 (TN, FP)
	Pos	2581	3888 (FN, TP)

Metrics:

Accuracy:	0.5997
Precision:	0.4199
Recall:	0.6010
F1-Score:	0.4944
Specificity:	0.5991
ROC AUC:	0.6411
PR AUC:	0.4491

```
>>> XGBoost complete.
```

```
In [35]: # =====
# 6E. ADVANCED MODEL: LightGBM
# =====
print("\n" + "="*60)
print(">>> STEP 6E: Training LightGBM...")
print("=".*60)

lgb_model = LGBMClassifier(
```

```

    n_estimators=200,
    max_depth=6,
    learning_rate=0.05,
    subsample=0.8,
    colsample_bytree=0.8,
    class_weight="balanced",
    random_state=RANDOM_SEED,
    n_jobs=-1,
    verbose=-1
)

lgb_model.fit(X_train, y_train)

y_pred_lgb = lgb_model.predict(X_test)
y_proba_lgb = lgb_model.predict_proba(X_test)[:, 1]

all_results["LightGBM"] = evaluate_model(
    "LightGBM",
    y_test, y_pred_lgb, y_proba_lgb
)

trained_models["LightGBM"] = lgb_model

print("\n>>> LightGBM complete.")

```

=====

>>> STEP 6E: Training LightGBM...

=====

=====

LightGBM

=====

Confusion Matrix:

		Predicted		
		Neg	Pos	
Actual	Neg	8077	5322	(TN, FP)
	Pos	2591	3878	(FN, TP)

Metrics:

Accuracy:	0.6017
Precision:	0.4215
Recall:	0.5995
F1-Score:	0.4950
Specificity:	0.6028
ROC AUC:	0.6429
PR AUC:	0.4516

>>> LightGBM complete.

In [36]:

```

# =====
# 6F. BASELINE MODEL: Logistic Regression
# =====
print("\n" + "="*60)
print(">>> STEP 6F: Training Logistic Regression...")
print("=".*60)

lr_model = LogisticRegression(
    class_weight="balanced",
    solver="liblinear",

```

```

        max_iter=1000,
        random_state=RANDOM_SEED
    )

lr_model.fit(X_train, y_train)

y_pred_lr = lr_model.predict(X_test)
y_proba_lr = lr_model.predict_proba(X_test)[:, 1]

all_results["LogisticRegression"] = evaluate_model(
    "Logistic Regression",
    y_test, y_pred_lr, y_proba_lr
)

trained_models["LogisticRegression"] = lr_model

print("\n>>> Logistic Regression complete.")

```

=====
>>> STEP 6F: Training Logistic Regression...
=====

=====
Logistic Regression
=====

Confusion Matrix:

		Predicted		
		Neg	Pos	
Actual	Neg	7312	6087	(TN, FP)
	Pos	2311	4158	(FN, TP)

Metrics:

Accuracy:	0.5773
Precision:	0.4059
Recall:	0.6428
F1-Score:	0.4975
Specificity:	0.5457
ROC AUC:	0.6249
PR AUC:	0.4149

>>> Logistic Regression complete.

In [37]: # ======
6G. MODEL: K-Nearest Neighbors (KNN)
======
print("\n" + "="*60)
print(">>> STEP 6G: Training KNN...")
print("=*60)

KNN requires feature scaling
knn_pipeline = Pipeline([
 ("scaler", StandardScaler()),
 ("knn", KNeighborsClassifier(n_neighbors=11, n_jobs=-1))
])

knn_pipeline.fit(X_train, y_train)

y_pred_knn = knn_pipeline.predict(X_test)
y_proba_knn = knn_pipeline.predict_proba(X_test)[:, 1]

```

all_results["KNN"] = evaluate_model(
    "K-Nearest Neighbors (k=11)",
    y_test, y_pred_knn, y_proba_knn
)

trained_models["KNN"] = knn_pipeline

print("\n>>> KNN complete.")

```

=====

>>> STEP 6G: Training KNN...

=====

=====

K-Nearest Neighbors (k=11)

=====

Confusion Matrix:

		Predicted		
		Neg	Pos	
Actual	Neg	11992	1407	(TN, FP)
	Pos	5147	1322	(FN, TP)

Metrics:

Accuracy:	0.6701
Precision:	0.4844
Recall:	0.2044
F1-Score:	0.2875
Specificity:	0.8950
ROC AUC:	0.5902
PR AUC:	0.3964

>>> KNN complete.

```

In [38]: # =====
# 6I. MODEL: Naive Bayes
# =====
print("\n" + "="*60)
print(">>> STEP 6I: Training Naive Bayes...")
print("=".*60)

nb_model = GaussianNB()

nb_model.fit(X_train, y_train)

y_pred_nb = nb_model.predict(X_test)
y_proba_nb = nb_model.predict_proba(X_test)[:, 1]

all_results["NaiveBayes"] = evaluate_model(
    "Naive Bayes",
    y_test, y_pred_nb, y_proba_nb
)

trained_models["NaiveBayes"] = nb_model

print("\n>>> Naive Bayes complete.")

```

```
=====
>>> STEP 6I: Training Naive Bayes...
=====

=====
Naive Bayes
=====

Confusion Matrix:
      Predicted
      Neg   Pos
Actual Neg  10322  3077  (TN, FP)
Actual Pos   4078  2391  (FN, TP)

Metrics:
  Accuracy:  0.6399
  Precision: 0.4373
  Recall: 0.3696
  F1-Score: 0.4006
  Specificity: 0.7704
  ROC AUC: 0.6155
  PR AUC: 0.4103

>>> Naive Bayes complete.
```

In [39]:

```
# =====
# 6J. 10-FOLD CROSS-VALIDATION (ALL MODELS EXCEPT SVM)
# =====
print("\n" + "="*60)
print("=>> STEP 6J: 10-Fold Cross-Validation (Paper-style Evaluation)")
print("=*60)

# Models to evaluate with 10-fold CV
# Note: SVM excluded due to extremely long training time
models_for_cv = {
    "RF_Baseline": RandomForestClassifier(
        n_estimators=300, max_depth=10, min_samples_leaf=2,
        class_weight="balanced", random_state=RANDOM_SEED, n_jobs=-1
    ),
    "RF_Tuned": RandomForestClassifier(
        **grid_search_rf.best_params_,
        class_weight="balanced", random_state=RANDOM_SEED, n_jobs=-1
    ),
    "XGBoost": XGBClassifier(
        n_estimators=200, max_depth=6, learning_rate=0.1,
        subsample=0.8, colsample_bytree=0.8,
        scale_pos_weight=scale_pos_weight,
        random_state=RANDOM_SEED, n_jobs=-1, verbosity=0
    ),
    "LightGBM": LGBMClassifier(
        n_estimators=200, max_depth=6, learning_rate=0.05,
        subsample=0.8, colsample_bytree=0.8,
        class_weight="balanced", random_state=RANDOM_SEED, n_jobs=-1, verbose=-1
    ),
    "LogisticRegression": LogisticRegression(
        class_weight="balanced", solver="liblinear",
        max_iter=1000, random_state=RANDOM_SEED
    ),
    "KNN": Pipeline([
        ("scaler", StandardScaler())
    ])
}
```

```

        ("knn", KNeighborsClassifier(n_neighbors=11, n_jobs=-1))
    ]),
    # SVM excluded: too slow for 10-fold CV on this dataset
    "NaiveBayes": GaussianNB()
}

# Store CV results
cv_results = {}

print("\nRunning 10-fold CV for all models (SVM excluded)...")

for name, model in models_for_cv.items():
    print(f" Evaluating {name}...", end=" ")

    # Accuracy
    acc_scores = cross_val_score(model, X, y, cv=cv_10fold, scoring="accuracy",

    # F1-score
    f1_scores = cross_val_score(model, X, y, cv=cv_10fold, scoring="f1", n_jobs=1)

    # Precision
    prec_scores = cross_val_score(model, X, y, cv=cv_10fold, scoring="precision"

    # Recall
    rec_scores = cross_val_score(model, X, y, cv=cv_10fold, scoring="recall", n_jobs=1)

    # ROC AUC
    auc_scores = cross_val_score(model, X, y, cv=cv_10fold, scoring="roc_auc", n_jobs=1)

    cv_results[name] = {
        "accuracy_mean": acc_scores.mean(),
        "accuracy_std": acc_scores.std(),
        "f1_mean": f1_scores.mean(),
        "f1_std": f1_scores.std(),
        "precision_mean": prec_scores.mean(),
        "precision_std": prec_scores.std(),
        "recall_mean": rec_scores.mean(),
        "recall_std": rec_scores.std(),
        "roc_auc_mean": auc_scores.mean(),
        "roc_auc_std": auc_scores.std()
    }

    print(f"Done. F1={f1_scores.mean():.4f} ({±{f1_scores.std():.4f}})")

# Add SVM results from test set only (no CV)
# Mark it separately so we know it's not from 10-fold CV
print("\n Note: SVM evaluated on test set only (10-fold CV too slow)")

# =====
# PRINT CV RESULTS TABLE
# =====
print("\n" + "="*90)
print(" 10-FOLD CROSS-VALIDATION RESULTS SUMMARY")
print("="*90)
print(f"{'Model':<20} {'Accuracy':>14} {'Precision':>14} {'Recall':>14} {'F1-Score':>14}")
print("-"*90)

for name, res in cv_results.items():
    print(f"{name:<20} "
          f"{res['accuracy_mean']:.4f}±{res['accuracy_std']:.4f} "

```

```
f" {res['precision_mean']:.4f}±{res['precision_std']:.4f} "
f" {res['recall_mean']:.4f}±{res['recall_std']:.4f} "
f" {res['f1_mean']:.4f}±{res['f1_std']:.4f} "
f" {res['roc_auc_mean']:.4f}±{res['roc_auc_std']:.4f}" )

print("-"*90)
print("Note: SVM excluded from 10-fold CV due to computational time constraints.")
print("      SVM results available from test set evaluation (Step 6H).")
print("=*90)

# Find best model by F1
best_model_name = max(cv_results, key=lambda x: cv_results[x]["f1_mean"])
print(f"\nBest model by F1-score: {best_model_name} (F1 = {cv_results[best_model_name]['f1_mean']}")

print("\n>>> 10-Fold Cross-Validation complete.")
```

```
=====
>>> STEP 6J: 10-Fold Cross-Validation (Paper-style Evaluation)
=====
```

Running 10-fold CV for all models (SVM excluded)...

Done. F1=0.5018 (± 0.0053)..
 Done. F1=0.5021 (± 0.0058)
 Done. F1=0.5001 (± 0.0051)
 Done. F1=0.5009 (± 0.0050)
 Done. F1=0.5036 (± 0.0048)
 session...
 Done. F1=0.3974 (± 0.0161)
 Done. F1=0.3956 (± 0.0269)..

Note: SVM evaluated on test set only (10-fold CV too slow)

```
=====
=====
10-FOLD CROSS-VALIDATION RESULTS SUMMARY
=====
```

Model ROC AUC	Accuracy	Precision	Recall	F1-Score
<hr/>				
RF_Baseline 79±0.0071	0.6054±0.0029	0.4260±0.0033	0.6105±0.0100	0.5018±0.0053 0.64
RF_Tuned 51±0.0073	0.5997±0.0048	0.4220±0.0041	0.6201±0.0147	0.5021±0.0058 0.64
XGBoost 73±0.0070	0.6079±0.0045	0.4276±0.0044	0.6024±0.0088	0.5001±0.0051 0.64
LightGBM 84±0.0071	0.6090±0.0043	0.4286±0.0042	0.6028±0.0096	0.5009±0.0050 0.64
LogisticRegression 24±0.0068	0.5818±0.0037	0.4104±0.0036	0.6515±0.0077	0.5036±0.0048 0.63
KNN 40±0.0095	0.6384±0.0165	0.4374±0.0242	0.3674±0.0338	0.3974±0.0161 0.60
NaiveBayes 29±0.0073	0.6437±0.0060	0.4419±0.0101	0.3600±0.0408	0.3956±0.0269 0.62
<hr/>				

Note: SVM excluded from 10-fold CV due to computational time constraints.
 SVM results available from test set evaluation (Step 6H).

Best model by F1-score: LogisticRegression (F1 = 0.5036)

>>> 10-Fold Cross-Validation complete.

```
In [40]: # =====
# 7. COMPREHENSIVE MODEL EVALUATION
# =====
print("\n" + "="*60)
print(">>> STEP 7: Comprehensive Model Evaluation")
print("=".*60)

# =====
# 7A. TEST SET RESULTS SUMMARY TABLE
# =====
```

```

print("\n>>> STEP 7A: Test Set Evaluation Summary")

# List of models to summarize (excluding SVM)
models_to_summarize = ["RF_Baseline", "RF_Tuned", "XGBoost", "LightGBM",
                      "LogisticRegression", "KNN", "NaiveBayes"]

print("\n" + "="*100)
print(" TEST SET EVALUATION RESULTS (All Models)")
print("=*100")
print(f"{'Model':<20} {'Accuracy':>10} {'Precision':>10} {'Recall':>10} {'F1-Sco
print("-"*100)

for name in models_to_summarize:
    if name in all_results:
        res = all_results[name]
        print(f"{name:<20} "
              f"{res['accuracy']):>10.4f} "
              f"{res['precision']):>10.4f} "
              f"{res['recall']):>10.4f} "
              f"{res['f1_score']):>10.4f} "
              f"{res['specificity']):>12.4f} "
              f"{res['roc_auc']):>10.4f} "
              f"{res['pr_auc']):>10.4f}")

print("=*100")

# Confusion Matrix Summary
print("\n" + "="*70)
print(" CONFUSION MATRIX SUMMARY (TN, FP, FN, TP)")
print("=*70")
print(f"{'Model':<20} {'TN':>10} {'FP':>10} {'FN':>10} {'TP':>10}")
print("-"*70)

for name in models_to_summarize:
    if name in all_results:
        res = all_results[name]
        print(f"{name:<20} "
              f"{res['tn']):>10} "
              f"{res['fp']):>10} "
              f"{res['fn']):>10} "
              f"{res['tp']):>10}")

print("=*70")

# Find best model by F1 on test set
best_test_model = max(models_to_summarize, key=lambda x: all_results[x]["f1_sc
print(f"\nBest model on test set by F1-score: {best_test_model} (F1 = {all_resul

```

```
=====
>>> STEP 7: Comprehensive Model Evaluation
=====
```

```
>>> STEP 7A: Test Set Evaluation Summary
=====
```

```
=====
===== TEST SET EVALUATION RESULTS (All Models)
=====
```

Model		Accuracy	Precision	Recall	F1-Score	Specificity
ROC AUC	PR AUC					
<hr/>						
RF_Baseline		0.5991	0.4199	0.6061	0.4961	0.5956
0.6418	0.4496					
RF_Tuned		0.5937	0.4163	0.6166	0.4970	0.5826
0.6377	0.4451					
XGBoost		0.5997	0.4199	0.6010	0.4944	0.5991
0.6411	0.4491					
LightGBM		0.6017	0.4215	0.5995	0.4950	0.6028
0.6429	0.4516					
LogisticRegression		0.5773	0.4059	0.6428	0.4975	0.5457
0.6249	0.4149					
KNN		0.6701	0.4844	0.2044	0.2875	0.8950
0.5902	0.3964					
NaiveBayes		0.6399	0.4373	0.3696	0.4006	0.7704
0.6155	0.4103					
<hr/>						

```
=====
===== CONFUSION MATRIX SUMMARY (TN, FP, FN, TP)
=====
```

Model	TN	FP	FN	TP
<hr/>				
RF_Baseline	7981	5418	2548	3921
RF_Tuned	7806	5593	2480	3989
XGBoost	8027	5372	2581	3888
LightGBM	8077	5322	2591	3878
LogisticRegression	7312	6087	2311	4158
KNN	11992	1407	5147	1322
NaiveBayes	10322	3077	4078	2391
<hr/>				

Best model on test set by F1-score: LogisticRegression (F1 = 0.4975)

```
In [41]: # =====
# 7B. ROC CURVES (ALL MODELS)
# =====
print("\n>>> STEP 7B: Plotting ROC Curves for All Models...")

plt.figure(figsize=(10, 8))

# Define colors for each model
colors = {
    "RF_Baseline": "#1f77b4",
    "RF_Tuned": "#ff7f0e",
    "XGBoost": "#2ca02c",
```

```

    "LightGBM": "#d62728",
    "LogisticRegression": "#9467bd",
    "KNN": "#8c564b",
    "NaiveBayes": "#e377c2"
}

# Plot ROC curve for each model
for name in models_to_summarize:
    if name in all_results:
        res = all_results[name]
        fpr, tpr, _ = roc_curve(y_test, res["y_proba"])
        auc_score = res["roc_auc"]
        plt.plot(fpr, tpr, color=colors[name], lw=2,
                  label=f"{name} (AUC = {auc_score:.4f})")

# Plot diagonal line (random classifier)
plt.plot([0, 1], [0, 1], color="gray", lw=2, linestyle="--", label="Random (AUC = 0.5)")


plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate (1 - Specificity)", fontsize=12)
plt.ylabel("True Positive Rate (Sensitivity/Recall)", fontsize=12)
plt.title("ROC Curves - All Models Comparison", fontsize=14)
plt.legend(loc="lower right", fontsize=10)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig("roc_curves_all_models.png", dpi=300, bbox_inches="tight")
plt.show()

print("ROC curves saved to 'roc_curves_all_models.png'")

```

>>> STEP 7B: Plotting ROC Curves for All Models...

ROC Curves - All Models Comparison



```
ROC curves saved to 'roc_curves_all_models.png'
```

```
In [42]: # =====
# 7C. PRECISION-RECALL CURVES (ALL MODELS)
# =====
print("\n>>> STEP 7C: Plotting Precision-Recall Curves for All Models...")

plt.figure(figsize=(10, 8))


# Plot PR curve for each model
for name in models_to_summarize:
    if name in all_results:
        res = all_results[name]
        precision, recall, _ = precision_recall_curve(y_test, res["y_proba"])
        pr_auc = res["pr_auc"]
        plt.plot(recall, precision, color=colors[name], lw=2,
                  label=f"{name} (PR AUC = {pr_auc:.4f})")

# Plot baseline (proportion of positive class)
baseline = y_test.mean()
plt.axhline(y=baseline, color="gray", lw=2, linestyle="--",
            label=f"Baseline (Prevalence = {baseline:.4f})")

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("Recall (Sensitivity)", fontsize=12)
plt.ylabel("Precision", fontsize=12)
plt.title("Precision-Recall Curves - All Models Comparison", fontsize=14)
plt.legend(loc="upper right", fontsize=10)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.savefig("pr_curves_all_models.png", dpi=300, bbox_inches="tight")
plt.show()

print("Precision-Recall curves saved to 'pr_curves_all_models.png'")
```

```
>>> STEP 7C: Plotting Precision-Recall Curves for All Models...
```



Precision-Recall curves saved to 'pr_curves_all_models.png'

```
In [43]: # =====#
# 7D. SHAP ANALYSIS (TREE-BASED MODELS)
# =====#
print("\n>>> STEP 7D: SHAP Analysis for Tree-based Models...")

# Initialize SHAP
shap.initjs()

# Sample data for SHAP (to speed up computation)
if len(X_test) > 1000:
    shap_sample_idx = np.random.choice(len(X_test), size=1000, replace=False)
    X_shap = X_test.iloc[shap_sample_idx]
else:
    X_shap = X_test

print(f"Using {len(X_shap)} samples for SHAP analysis")

# Models to analyze with SHAP (tree-based models only)
shap_models = {
    "RF_Tuned": trained_models["RF_Tuned"],
    "XGBoost": trained_models["XGBoost"],
    "LightGBM": trained_models["LightGBM"]
}

# Store SHAP values for each model
shap_values_dict = {}

for name, model in shap_models.items():
    print(f"\n>>> Computing SHAP values for {name}...")
```

```
# Create explainer
explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_shap)

# Handle different SHAP output formats
if isinstance(shap_values, list):
    # For RF: shap_values is [class_0, class_1], we want class_1
    shap_values_to_plot = shap_values[1]
else:
    shap_values_to_plot = shap_values

shap_values_dict[name] = shap_values_to_plot

# Plot SHAP summary
print(f"\n--- SHAP Summary Plot: {name} ---")
plt.figure(figsize=(10, 8))
shap.summary_plot(shap_values_to_plot, X_shap, feature_names=feature_cols, s
plt.title(f"SHAP Summary Plot - {name}", fontsize=14)
plt.tight_layout()
plt.savefig(f"shap_summary_{name}.png", dpi=300, bbox_inches="tight")
plt.show()
print(f"Saved to 'shap_summary_{name}'.png'")

print("\n>>> SHAP Summary Plots complete.")
```


>>> STEP 7D: SHAP Analysis for Tree-based Models...



Using 1000 samples for SHAP analysis

>>> Computing SHAP values for RF_Tuned...


--- SHAP Summary Plot: RF_Tuned ---
<Figure size 1000x800 with 0 Axes>



Saved to 'shap_summary_RF_Tuned.png'

>>> Computing SHAP values for XGBoost...


--- SHAP Summary Plot: XGBoost ---



Saved to 'shap_summary_XGBoost.png'

>>> Computing SHAP values for LightGBM...

--- SHAP Summary Plot: LightGBM ---



Saved to 'shap_summary_LightGBM.png'

>>> SHAP Summary Plots complete.

```

In [44]: # =====
# 7D-2. SHAP FEATURE IMPORTANCE (NUMERICAL)
# =====
print("\n>>> STEP 7D-2: SHAP Feature Importance (Numerical)...")


# Store feature importance for each model
feature_importance_dict = {}

print("\n" + "="*80)
print(" SHAP GLOBAL FEATURE IMPORTANCE (Mean |SHAP Value|)")
print("=*80")

for name, shap_vals in shap_values_dict.items():
    # Compute mean absolute SHAP values
    mean_abs_shap = np.abs(shap_vals).mean(axis=0)

    # Ensure it's a 1D array
    mean_abs_shap = np.ravel(mean_abs_shap)

    # Handle Length mismatch
    n_features = min(len(feature_cols), len(mean_abs_shap))

    # Create feature importance list
    importance_list = [(feature_cols[i], float(mean_abs_shap[i])) for i in range(n_features)]
    importance_list_sorted = sorted(importance_list, key=lambda x: x[1], reverse=True)

    feature_importance_dict[name] = importance_list_sorted

    print(f"\n--- {name} ---")
    print(f"{'Rank':<6} {'Feature':<25} {'Mean |SHAP|':>12}")
    print("-"*45)
    for rank, (feat, val) in enumerate(importance_list_sorted, 1):
        print(f"{rank:<6} {feat:<25} {val:>12.6f}")

print("\n" + "="*80)

# =====
# 7D-3. COMBINED FEATURE IMPORTANCE BAR PLOT
# =====
print("\n>>> STEP 7D-3: Combined Feature Importance Bar Plot...")


fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for idx, (name, importance_list) in enumerate(feature_importance_dict.items()):
    ax = axes[idx]

    # Get top 10 features
    top_n = min(10, len(importance_list))
    features = [x[0] for x in importance_list[:top_n]][::-1]
    values = [x[1] for x in importance_list[:top_n]][::-1]

    ax.barh(range(top_n), values, color=colors[name])
    ax.set_yticks(range(top_n))
    ax.set_yticklabels(features)
    ax.set_xlabel("Mean |SHAP Value|")
    ax.set_title(f"{name}")

```

```
ax.grid(True, alpha=0.3, axis='x')

plt.suptitle("Feature Importance Comparison (SHAP)", fontsize=14, fontweight='bold')
plt.tight_layout()
plt.savefig("shap_feature_importance_comparison.png", dpi=300, bbox_inches="tight")
plt.show()

print("Saved to 'shap_feature_importance_comparison.png'")
```

>>> STEP 7D-2: SHAP Feature Importance (Numerical)...


=====			
SHAP GLOBAL FEATURE IMPORTANCE (Mean SHAP Value)			
=====			
--- RF_Tuned ---			
Rank	Feature	Mean	SHAP

1	blood_pressure	0.048308	
2	cholesterol	0.048308	
3	gender	0.013761	
4	race	0.013761	
5	uric_acid	0.011575	
6	pregnancy	0.011575	
7	heart_disease	0.005604	
8	obesity	0.005604	
9	blurred_vision	0.003251	
10	alcohol	0.000702	
11	age	0.000702	
--- XGBoost ---			
Rank	Feature	Mean	SHAP

1	heart_disease	0.326678	
2	age	0.261835	
3	blood_pressure	0.079052	
4	race	0.073691	
5	alcohol	0.045456	
6	cholesterol	0.036351	
7	gender	0.024546	
8	obesity	0.012587	
9	uric_acid	0.010532	
10	pregnancy	0.007771	
11	blurred_vision	0.000473	
--- LightGBM ---			
Rank	Feature	Mean	SHAP

1	heart_disease	0.338956	
2	age	0.258432	
3	race	0.070365	
4	blood_pressure	0.069635	
5	alcohol	0.046756	
6	cholesterol	0.033201	
7	gender	0.018722	
8	uric_acid	0.010836	
9	obesity	0.010803	
10	pregnancy	0.006644	
11	blurred_vision	0.000491	
=====			

>>> STEP 7D-3: Combined Feature Importance Bar Plot...



Saved to 'shap_feature_importance_comparison.png'

```
In [49]: from sklearn.metrics import f1_score, precision_score, recall_score, confusion_m

# =====
# 7E. COMPREHENSIVE THRESHOLD TUNING
# =====
print("\n" + "="*80)
print(">> STEP 7E: Threshold Optimization for All Candidate Models")
print("=*80")
print("Rationale: Given the class imbalance and the comparable performance of to")
print("we optimize the decision threshold for all candidates to maximize the F1-")
print("-" * 80)

# Select models to tune (Top performers + Baseline)
models_to_tune = ["LightGBM", "XGBoost", "RF_Tuned", "LogisticRegression"]
tuning_results = []

# Header for the output table
header = f"{'Model':<20} {'Default F1':<12} {'Optimized F1':<14} {'Best Threshold':<12}"
print(header)
print("-" * 80)

for name in models_to_tune:
    if name in all_results:
        # Get probability scores for the positive class
        y_proba = all_results[name]["y_proba"]
        y_true = y_test

        best_f1 = -1
        best_thr = -1
        best_prec = -1
        best_rec = -1

        # Grid search for the best threshold (0.10 to 0.90)
        thresholds = np.linspace(0.1, 0.9, 81)

        for thr in thresholds:
            y_pred_thr = (y_proba >= thr).astype(int)
            # Calculate F1 score (zero_division=0 handles cases with no positive
            current_f1 = f1_score(y_true, y_pred_thr, zero_division=0)

            if current_f1 > best_f1:
                best_f1 = current_f1
                best_thr = thr
                best_prec = precision_score(y_true, y_pred_thr, zero_division=0)
                best_rec = recall_score(y_true, y_pred_thr, zero_division=0)
```

```

# Calculate improvement over default threshold (0.5)
default_f1 = all_results[name]["f1_score"]
improvement = best_f1 - default_f1

# Print row in the table
print(f"{name:<20} {default_f1:<12.4f} {best_f1:<14.4f} {best_thr:<16.3f}

# Store detailed results for plotting or further analysis
tuning_results.append({
    "Model": name,
    "Default_F1": default_f1,
    "Optimized_F1": best_f1,
    "Best_Threshold": best_thr,
    "Precision": best_prec,
    "Recall": best_rec
})

print("-" * 80)
print("Observation: Threshold tuning significantly impacts models like Logistic")
print("demonstrating their potential competitiveness when properly calibrated.")

```

```

=====
>>> STEP 7E: Threshold Optimization for All Candidate Models
=====
Rationale: Given the class imbalance and the comparable performance of top model
s,
we optimize the decision threshold for all candidates to maximize the F1-score.

-----
Model          Default F1   Optimized F1   Best Threshold   Improvement
-----
LightGBM        0.4950      0.5177       0.410           +0.0227
XGBoost         0.4944      0.5168       0.420           +0.0224
RF_Tuned        0.4970      0.5135       0.430           +0.0164
LogisticRegression 0.4975      0.5110       0.370           +0.0135
-----
Observation: Threshold tuning significantly impacts models like Logistic Regressi
on,
demonstrating their potential competitiveness when properly calibrated.

```

In [46]:

```

# =====
# 7E-5. EXECUTION TIME COMPARISON
# =====
print("\n" + "="*60)
print(">>> EXECUTION TIME COMPARISON")
print("=*60)

import time

execution_times = {}

models_for_timing = {
    "RF_Tuned": rf_tuned,
    "XGBoost": xgb_model,
    "LightGBM": lgb_model,
    "LogisticRegression": lr_model,
    "KNN": knn_pipeline,
    "NaiveBayes": nb_model,
}

print("\nMeasuring 5-fold CV execution time for each model...\n")

```

```

for name, model in models_for_timing.items():
    print(f"  Timing {name}...", end=" ", flush=True)
    start_time = time.time()
    _ = cross_val_score(model, X_train, y_train, cv=5, scoring='f1', n_jobs=-1)
    elapsed = time.time() - start_time
    execution_times[name] = elapsed
    print(f"{elapsed:.2f} seconds")

# Plot
print("\n" + "*50")
print("EXECUTION TIME SUMMARY (5-fold CV)")
print("*50")
print(f"{'Model':<25} {'Time (seconds)':>15}")
print("-*50")
for name, t in sorted(execution_times.items(), key=lambda x: x[1]):
    print(f"{name:<25} {t:>15.2f}")
print("*50")

# Visualization
plt.figure(figsize=(10, 6))
models_sorted = sorted(execution_times.keys(), key=lambda x: execution_times[x])
times_sorted = [execution_times[m] for m in models_sorted]
plt.barh(models_sorted, times_sorted, color='steelblue')
plt.xlabel('Execution Time (seconds)', fontsize=12)
plt.title('Model Training Time Comparison (5-fold CV)', fontsize=14)
for i, t in enumerate(times_sorted):
    plt.text(t + 0.1, i, f'{t:.2f}s', va='center')
plt.tight_layout()
plt.savefig('execution_time_comparison.png', dpi=300, bbox_inches='tight')
plt.show()
print("Saved: execution_time_comparison.png")

```

```
=====
>>> EXECUTION TIME COMPARISON
=====
```


Measuring 5-fold CV execution time for each model...

1.56 secondsTuned...
1.93 secondsost...
1.77 secondshtGBM...
0.32 secondsisticRegression...
25.72 seconds..
0.22 secondsveBayes...

```
=====
EXECUTION TIME SUMMARY (5-fold CV)
=====
```

Model	Time (seconds)
NaiveBayes	0.22
LogisticRegression	0.32
RF_Tuned	1.56
LightGBM	1.77
XGBoost	1.93
KNN	25.72

Model Training Time Comparison (5-fold CV)



Saved: execution_time_comparison.png

```
In [47]: # =====
# 7F. FINAL MODEL SUMMARY
# =====
print("\n" + "="*80)
print(">>> STEP 7F: FINAL MODEL SUMMARY")
print("=*80)

# =====
# 7F-1. TEST SET RESULTS COMPARISON
# =====
print("\n" + "-"*80)
print(" TEST SET RESULTS COMPARISON")
print("-"*80)

# Create summary dataframe
summary_data = []
for name in models_to_summarize:
    if name in all_results:
        res = all_results[name]
        summary_data.append({
            "Model": name,
            "Accuracy": res["accuracy"],
            "Precision": res["precision"],
            "Recall": res["recall"],
            "F1-Score": res["f1_score"],
            "Specificity": res["specificity"],
            "ROC AUC": res["roc_auc"],
            "PR AUC": res["pr_auc"]
        })
summary_df = pd.DataFrame(summary_data)
summary_df = summary_df.sort_values("F1-Score", ascending=False).reset_index(drop=True)

print("\nRanked by F1-Score (Test Set):")
print(summary_df.to_string(index=False))

# =====
```

```

# 7F-2. 10-FOLD CV RESULTS COMPARISON
# =====
print("\n" + "-"*80)
print(" 10-FOLD CROSS-VALIDATION RESULTS COMPARISON")
print("-"*80)

cv_summary_data = []
for name, res in cv_results.items():
    cv_summary_data.append({
        "Model": name,
        "Accuracy": f"{res['accuracy_mean']:.4f}±{res['accuracy_std']:.4f}",
        "Precision": f"{res['precision_mean']:.4f}±{res['precision_std']:.4f}",
        "Recall": f"{res['recall_mean']:.4f}±{res['recall_std']:.4f}",
        "F1-Score": f"{res['f1_mean']:.4f}±{res['f1_std']:.4f}",
        "ROC AUC": f"{res['roc_auc_mean']:.4f}±{res['roc_auc_std']:.4f}",
        "F1_mean": res['f1_mean'] # for sorting
    })

cv_summary_df = pd.DataFrame(cv_summary_data)
cv_summary_df = cv_summary_df.sort_values("F1_mean", ascending=False).reset_index()
cv_summary_df = cv_summary_df.drop(columns=["F1_mean"])

print("\nRanked by F1-Score (10-Fold CV):")
print(cv_summary_df.to_string(index=False))

# =====
# 7F-3. BEST MODEL SUMMARY
# =====
print("\n" + "-"*80)
print("  BEST MODEL SUMMARY")
print("-"*80)

# Best model from 10-fold CV
best_cv_model = max(cv_results, key=lambda x: cv_results[x]["f1_mean"])
best_cv_f1 = cv_results[best_cv_model]["f1_mean"]

# Best model from test set
best_test_model = max(models_to_summarize, key=lambda x: all_results[x]["f1_score"])
best_test_f1 = all_results[best_test_model]["f1_score"]

print(f"\nBest Model (10-Fold CV): {best_cv_model} (F1 = {best_cv_f1:.4f})")
print(f"Best Model (Test Set): {best_test_model} (F1 = {best_test_f1:.4f})")
print(f"Optimal Threshold: {best_thr:.3f} (F1 = {best_f1:.4f})")

# =====
# 7F-4. KEY FINDINGS
# =====
print("\n" + "-"*80)
print("  KEY FINDINGS")
print("-"*80)

print("""
1. MODEL PERFORMANCE:
- At default threshold (0.5), all top models achieved similar F1 ≈ 0.50
- Logistic Regression had the highest default F1-score (0.5036 in 10-fold CV)
- After threshold optimization, LightGBM achieved the best F1-score (0.5177)
- Tree-based models (LightGBM, XGBoost, RF) and Logistic Regression performed well
- KNN and Naive Bayes performed poorly (F1 = 0.29 and 0.40, respectively)

2. CLASS IMBALANCE:
""")

```

```
- Positive class (diabetes2=1): 32.56%
- Negative class (diabetes2=0): 67.44%
- Imbalance ratio: 1:2.07
- All models used class_weight='balanced' or scale_pos_weight to address imba

3. FEATURE IMPORTANCE (SHAP):
- heart_disease and age are the dominant predictors in XGBoost and LightGBM
- race and blood_pressure showed moderate importance
- blurred_vision and pregnancy contributed minimally due to low prevalence

4. THRESHOLD TUNING:
- Default threshold (0.5) is suboptimal for this dataset
- Optimal threshold for LightGBM: 0.410
- F1 improvement from threshold tuning: 0.4950 -> 0.5177 (+4.6%)
- Lowering threshold increases recall (0.60 -> 0.81) at the cost of precision
"""
print("=*80")
print(">>> ANALYSIS COMPLETE")
print("=*80")
```

>>> STEP 7F: FINAL MODEL SUMMARY

TEST SET RESULTS COMPARISON

Ranked by F1-Score (Test Set):

Model	Accuracy	Precision	Recall	F1-Score	Specificity	ROC AUC
PR AUC						
LogisticRegression	0.577310 0.414935	0.405857 0.445133	0.642758 0.616633	0.497547 0.497041	0.545712 0.582581	0.624948 0.637662
RF_Tuned	0.593668 0.449574	0.416301 0.419852	0.616633 0.606122	0.497041 0.496078	0.595641 0.595641	0.641848 0.641848
RF_Baseline	0.599054 0.451610	0.419852 0.421522	0.606122 0.599474	0.496078 0.494990	0.602806 0.602806	0.642862 0.642862
LightGBM	0.601721 0.449085	0.421522 0.419870	0.599474 0.601020	0.494990 0.494373	0.602806 0.599075	0.641092 0.641092
XGBoost	0.599708 0.449085	0.419870 0.421522	0.601020 0.599474	0.494373 0.494990	0.602806 0.602806	0.641092 0.641092
NaiveBayes	0.639873 0.410307	0.437271 0.421522	0.369609 0.599474	0.400603 0.400603	0.770356 0.770356	0.615451 0.615451
KNN	0.670123 0.396366	0.484427 0.421522	0.204359 0.599474	0.287454 0.287454	0.894992 0.894992	0.590208 0.590208

10-FOLD CROSS-VALIDATION RESULTS COMPARISON

Ranked by F1-Score (10-Fold CV):

Model	Accuracy	Precision	Recall	F1-Score	
ROC AUC					
LogisticRegression	0.5818±0.0037 ±0.0068	0.4104±0.0036 0.4104±0.0036	0.6515±0.0077 0.6515±0.0077	0.5036±0.0048 0.5036±0.0048	0.6324 0.6324
RF_Tuned	0.5997±0.0048 ±0.0073	0.4220±0.0041 0.4220±0.0041	0.6201±0.0147 0.6201±0.0147	0.5021±0.0058 0.5021±0.0058	0.6451 0.6451
RF_Baseline	0.6054±0.0029 ±0.0071	0.4260±0.0033 0.4260±0.0033	0.6105±0.0100 0.6105±0.0100	0.5018±0.0053 0.5018±0.0053	0.6479 0.6479
LightGBM	0.6090±0.0043 ±0.0071	0.4286±0.0042 0.4286±0.0042	0.6028±0.0096 0.6028±0.0096	0.5009±0.0050 0.5009±0.0050	0.6484 0.6484
XGBoost	0.6079±0.0045 ±0.0070	0.4276±0.0044 0.4276±0.0044	0.6024±0.0088 0.6024±0.0088	0.5001±0.0051 0.5001±0.0051	0.6473 0.6473
KNN	0.6384±0.0165 ±0.0095	0.4374±0.0242 0.4374±0.0242	0.3674±0.0338 0.3674±0.0338	0.3974±0.0161 0.3974±0.0161	0.6040 0.6040
NaiveBayes	0.6437±0.0060 ±0.0073	0.4419±0.0101 0.4419±0.0101	0.3600±0.0408 0.3600±0.0408	0.3956±0.0269 0.3956±0.0269	0.6229 0.6229

BEST MODEL SUMMARY

Best Model (10-Fold CV): LogisticRegression (F1 = 0.5036)
Best Model (Test Set): LogisticRegression (F1 = 0.4975)
Optimal Threshold: 0.410 (F1 = 0.5177)

KEY FINDINGS

1. MODEL PERFORMANCE:

- Tree-based models (LightGBM, XGBoost, RF) performed best with $F1 \approx 0.32$
- LightGBM achieved the highest 10-fold CV F1-score (0.3216)
- Logistic Regression performed comparably with simpler interpretation
- KNN performed poorly on this dataset ($F1 \approx 0.10$)
- Naive Bayes had high recall (0.96) but low precision

2. CLASS IMBALANCE:

- Positive class (diabetes2=1): 16.26%
- Negative class (diabetes2=0): 83.74%
- All models used `class_weight='balanced'` to address imbalance

3. FEATURE IMPORTANCE (SHAP):

- Most important features vary slightly across models
- Clinical features derived from ICD-9 codes are key predictors

4. THRESHOLD TUNING:

- Default threshold (0.5) may not be optimal
- Optimal threshold found: 0.410
- F1 improvement from threshold tuning: 0.4950 -> 0.5177

```
=====
>>> ANALYSIS COMPLETE
=====
```

In [48]:

```
# =====
# PUBLICATION-READY SUMMARY TABLES
# =====
print("\n" + "="*80)
print(">>> PUBLICATION-READY SUMMARY TABLES")
print("=*80)

# Table 1: The Result of Test Set
print("\n" + "-"*80)
print("TABLE: Test Set Performance Metrics")
print("-"*80)

models_to_summarize = ["RF_Baseline", "RF_Tuned", "XGBoost", "LightGBM",
                      "LogisticRegression", "KNN", "NaiveBayes"]

table_data = []
for name in models_to_summarize:
    if name in all_results:
        r = all_results[name]
        table_data.append({
            'Model': name,
            'Accuracy': f'{r["accuracy"]:.4f}',
            'Precision': f'{r["precision"]:.4f}',
            'Recall': f'{r["recall"]:.4f}',
            'F1-Score': f'{r["f1_score"]:.4f}',
            'ROC-AUC': f'{r["roc_auc"]:.4f}',
            'PR-AUC': f'{r["pr_auc"]:.4f}'
        })
table_df = pd.DataFrame(table_data)
print(table_df.to_string(index=False))

# CSV
table_df.to_csv('model_performance_summary.csv', index=False)
print("\nSaved: model_performance_summary.csv")
```

```
=====
>>> PUBLICATION-READY SUMMARY TABLES
=====
```

TABLE: Test Set Performance Metrics

	Model	Accuracy	Precision	Recall	F1-Score	ROC-AUC	PR-AUC
	RF_Baseline	0.5991	0.4199	0.6061	0.4961	0.6418	0.4496
	RF_Tuned	0.5937	0.4163	0.6166	0.4970	0.6377	0.4451
	XGBoost	0.5997	0.4199	0.6010	0.4944	0.6411	0.4491
	LightGBM	0.6017	0.4215	0.5995	0.4950	0.6429	0.4516
	LogisticRegression	0.5773	0.4059	0.6428	0.4975	0.6249	0.4149
	KNN	0.6701	0.4844	0.2044	0.2875	0.5902	0.3964
	NaiveBayes	0.6399	0.4373	0.3696	0.4006	0.6155	0.4103

Saved: model_performance_summary.csv